

CPSC532W: HOMEWORK 1

Name: Adam Jozefiak, Student No.: 27458158

1. We show that Gamma distribution is conjugate to the Poisson distribution. Let

$$\begin{aligned}\lambda &\sim \text{Gamma}(\alpha, \beta) \\ x_1, \dots, x_N &\sim \text{Poisson}(\lambda)\end{aligned}$$

Then by Bayes rule, the posterior distribution $p(\lambda|\mathbf{x})$ can be written as

$$p(\lambda|x_{1:N}) = \frac{p(x_{1:N}|\lambda)p(\lambda)}{p(x_{1:N})} \propto p(x_{1:N}|\lambda)p(\lambda)$$

where the last proportionality follows by the fact that $p(x_{1:N}) = \int p(x_{1:N}|\lambda)p(\lambda)d\lambda$ is constant.

We note the probability density functions of $\text{Gamma}(\alpha, \beta)$ and $\text{Poisson}(\lambda)$:

$$p(\gamma) = \frac{\beta^\alpha}{\Gamma(\alpha)} \lambda^{\alpha-1} e^{-\beta\lambda}$$

$$p(x_{1:N}) = \prod_{i=1}^N \frac{\lambda^{x_i} e^{-\lambda}}{x_i!}$$

Then,

$$\begin{aligned}p(\lambda|x_{1:N}) &\propto p(x_{1:N}|\lambda)p(\lambda) && \text{as argued earlier} \\ &= \left(\prod_{i=1}^N \frac{\lambda^{x_i} e^{-\lambda}}{x_i!}\right) \left(\frac{\beta^\alpha}{\Gamma(\alpha)} \lambda^{\alpha-1} e^{-\beta\lambda}\right) \\ &= \frac{\beta^\alpha}{\prod_{i=1}^N x_i!} \lambda^{\sum_{i=1}^N x_i + \alpha - 1} e^{-(n+\beta)\lambda} \\ &\propto \lambda^{\sum_{i=1}^N x_i + \alpha - 1} e^{-(n+\beta)\lambda} && \text{by } \lambda \text{ not appearing in } \frac{\beta^\alpha}{\prod_{i=1}^N x_i!}\end{aligned}$$

Therefore, up to a normalizing constant, $p(\lambda|x_{1:N}) \propto \lambda^{\sum_{i=1}^N x_i + \alpha - 1} e^{-(n+\beta)\lambda}$ and hence $p(\lambda|x_{1:N}) \sim \text{Gamma}(\sum_{i=1}^N x_i + \alpha, n + \beta)$. Therefore, the a prior Gamma distribution is a conjugate prior for a likelihood Poisson distribution.

2. Let x and x' be vectors of random variables. Then the Gibbs transition operator, for the change in the k^{th} variable or index of x , is defined as $T_k(x, x') = p(x'_k|x_{-k}) \prod_{i \neq k} \mathbb{I}(x_i = x'_i)$. Then we can see that the Gibbs transition operator satisfies the detailed balance equation by the following argument

$$\begin{aligned}
& p(x)T(x, x') \\
&= p(x)p(x'_k|x_{-k}) \prod_{i \neq k} \mathbb{I}(x_i = x'_i) \\
&= p(x_k|x_{-k})p(x_{-k})p(x'_k|x_{-k}) \prod_{i \neq k} \mathbb{I}(x_i = x'_i) \quad \text{by the product rule} \\
&= p(x')p(x_k|x_{-k}) \prod_{i \neq k} \mathbb{I}(x_i = x'_i) \quad \text{by the product rule on } p(x') = p(x'_k|x_{-k})p(x_{-k}) \\
&= p(x')p(x_k|x'_{-k}) \prod_{i \neq k} \mathbb{I}(x_i = x'_i) \quad \text{by } p(x_k|x_{-k}) \prod_{i \neq k} \mathbb{I}(x_i = x'_i) > 0 \Rightarrow x'_{-k} = x_{-k} \\
&= p(x')T(x', x)
\end{aligned}$$

This completes the proof of the Gibbs transition operator satisfying the detailed balance equation and therefore Gibbs transition operator can be interpreted as a Metropolis-Hastings transition operator that always accepts.

3. My solution to this problem is in Julia, based off the sample code. I found that the probability of it being cloudy given that the grass is wet is 57.58% by enumerating all possible world states and conditioning by counting which proportion are cloudy given the grass is wet. Using ancestral sampling and rejection, with 10000 successful samples, I found that the probability that it is cloudy given the grass is wet is 57.50% while 35.43% of samples were rejected. Using Gibbs sampling I found that the probability that it is cloudy given that the grass is wet is 58.66%. In my implementation, an index of 2 indicates a “true” assignment and an index 1 indicates a “false” assignment to a random variable.

Below is my code for computing the probability by enumerating all possible world states and conditioning by counting which proportion are cloudy given the grass is wet:

```

45  ## condition and marginalize:
46  p_C_given_W = 0.0
47  p_C_and_W = 0.0
48  p_W = 0.0
49  for c in 1:2
50      for s in 1:2
51          for r in 1:2
52              global p_W += p[c,s,r,2]
53          end
54      end
55  end
56  for s in 1:2
57      for r in 1:2
58          global p_C_and_W += p[2,s,r,2]
59      end
60  end
61  p_C_given_W = p_C_and_W / p_W
62
63  println("There is a ", p_C_given_W*100, "% chance it is cloudy given the grass is wet.")

```

Below is my code for computing the probability by ancestral sampling and rejection:

```

65  ##2. ancestral sampling and rejection:
66  num_samples = 10000
67  samples = zeros(num_samples)
68  rejections = 0
69  i = 1
70  while i <= num_samples
71      C = 0
72      S = 0
73      R = 0
74      W = 0
75      p = rand()
76      if p > p_C(1)
77          C = 2
78      else
79          C = 1
80      end
81      p = rand()
82      if p > p_S_given_C(1,C)
83          S = 2
84      else
85          S = 1
86      end
87      p = rand()
88      if p > p_R_given_C(1,C)
89          R = 2
90      else
91          R = 1
92      end
93      p = rand()
94      if p > p_W_given_S_R(1,S,R)
95          W = 2
96      else
97          W = 1
98      end
99      if W == 2
100         global samples[i] = C-1
101         global i += 1
102     else
103         global rejections += 1
104     end
105 end
106
107
108 println("The chance of it being cloudy given the grass is wet is ", (sum(samples)./num_samples)*100, "%.")
109 println(100*rejections/(num_samples+rejections), "% of the total samples were rejected.")

```

Below is my code for computing the probability by Gibbs sampling:

```

139 ##gibbs sampling
140 # num_samples = 10000
141 num_samples = 10000
142 samples = zeros(num_samples)
143 state = ones(Int64,4)
144 #c,s,r,w, set w = True
145 state[4] = 2
146 i = 1
147 while i <= num_samples
148
149     # Sample r
150     prob = rand()
151     if prob > p_R_given_C_S_W[state[1],state[2],1,state[4]]
152         state[3] = 2
153     else
154         state[3] = 1
155     end
156
157     # Sample s
158     prob = rand()
159     if prob > p_S_given_C_R_W[state[1],1,state[3],state[4]]
160         state[2] = 2
161     else
162         state[2] = 1
163     end
164
165     # Sample c
166     prob = rand()
167     if prob > p_C_given_S_R[1,state[2],state[3]]
168         state[1] = 2
169     else
170         state[1] = 1
171     end
172     samples[i] = state[1]
173     global i += 1
174
175 end
176
177
178 println("The chance of it being cloudy given the grass is wet is, ",100*sum(samples .- 1)/num_samples ,"%.")

```

4. • We first derive the updates for Metropolis Hastings sampling of the \hat{t} block, that is we want to sample from $p(\hat{t}|\hat{x}, t, x, w, \sigma^2, \alpha)$.
 Let the proposal distribution be $\hat{t}'|\hat{t} \sim \text{Normal}(\hat{t}, \sigma^2)$. Then $\hat{t}'|\hat{t}$ and $\hat{t}|\hat{t}'$ are symmetric, i.e $q(\hat{t}'|\hat{t}) = q(\hat{t}|\hat{t}')$, and so $\frac{q(\hat{t}|\hat{t}')}{q(\hat{t}'|\hat{t})} = 1$. Then after sampling \hat{t}' from $q(\hat{t}'|\hat{t})$ the accepting probability is:

$$\begin{aligned}
& A(\hat{t}'|\hat{t}) \\
&= \min(1, \frac{p(\hat{t}'|\hat{x}, t, x, w, \sigma^2, \alpha)q(\hat{t}|\hat{t}')}{p(\hat{t}|\hat{x}, t, x, w, \sigma^2, \alpha)q(\hat{t}'|\hat{t})}) \\
&= \min(1, \frac{p(\hat{t}'|\hat{x}, t, x, w, \sigma^2, \alpha)}{p(\hat{t}|\hat{x}, t, x, w, \sigma^2, \alpha)}) \quad \text{by the symmetry of the proposal distribution} \\
&= \min(1, \frac{p(\hat{t}', \hat{x}, t, x, w, \sigma^2, \alpha)p(\hat{x}, t, x, w, \sigma^2, \alpha)}{p(\hat{t}, \hat{x}, t, x, w, \sigma^2, \alpha)p(\hat{x}, t, x, w, \sigma^2, \alpha)}) \quad \text{by the product rule} \\
&= \min(1, \frac{p(\hat{t}', \hat{x}, t, x, w, \sigma^2, \alpha)}{p(\hat{t}, \hat{x}, t, x, w, \sigma^2, \alpha)}) \\
&= \min(1, \frac{\prod_{i=1}^N p(t_i|x, w, \sigma^2)p(w|\alpha)p(\hat{t}'|\hat{x}, w, \sigma^2)}{\prod_{i=1}^N p(t_i|x, w, \sigma^2)p(w|\alpha)p(\hat{t}|\hat{x}, w, \sigma^2)}) \\
&= \min(1, \frac{p(\hat{t}'|\hat{x}, w, \sigma^2)}{p(\hat{t}|\hat{x}, w, \sigma^2)})
\end{aligned}$$

Therefore, in the Metropolis Hastings sampling block for \hat{t} , we sample \hat{t}' from $\text{Normal}(\hat{t}, \sigma^2)$ and we accept \hat{t}' with probability $\min(1, \frac{p(\hat{t}'|\hat{x}, w, \sigma^2)}{p(\hat{t}|\hat{x}, w, \sigma^2)})$.

Next, we derive the updates for Metropolis Hastings sampling of the w block, that is we want to sample from $p(w|\hat{t}, \hat{x}, t, x, \sigma^2, \alpha)$. We can treat the \hat{t} and \hat{x} as t and x and so we will derive the updates for Metropolis Hastings sampling of $p(w|t, x, \sigma^2, \alpha)$.

Let the proposal distribution be $w'|w \sim \text{Normal}(w, \alpha I)$. Then $w'|w$ and $w|w'$ are symmetric, i.e $q(w'|w) = q(w|w')$, and so $\frac{q(w|w')}{q(w'|w)} = 1$. Then after sampling w' from $q(w'|w)$ the accepting probability is:

$$\begin{aligned}
& A(w'|w) \\
&= \min(1, \frac{p(w'|t, x, \sigma^2, \alpha)q(w|w')}{p(w|t, x, \sigma^2, \alpha)q(w'|w)}) \\
&= \min(1, \frac{p(w'|t, x, \sigma^2, \alpha)}{p(w|t, x, \sigma^2, \alpha)}) \quad \text{by the symmetry of the proposal distribution} \\
&= \min(1, \frac{p(w', t, x, \sigma^2, \alpha)p(t, x, \sigma^2, \alpha)}{p(w, t, x, \sigma^2, \alpha)p(t, x, \sigma^2, \alpha)}) \quad \text{by the product rule} \\
&= \min(1, \frac{p(w', t, x, \sigma^2, \alpha)}{p(w, t, x, \sigma^2, \alpha)}) \\
&= \min(1, \frac{\prod_{i=1}^N p(t_i|x, w, \sigma^2)p(w|\alpha)}{\prod_{i=1}^N p(t_i|x, w', \sigma^2)p(w'|\alpha)})
\end{aligned}$$

Therefore, in the Metropolis Hastings sampling block for w , we sample w' from $\text{Normal}(w, \alpha I)$ and we accept w' with probability $\min(1, \frac{\prod_{i=1}^N p(t_i|x, w, \sigma^2)p(w|\alpha)}{\prod_{i=1}^N p(t_i|x, w', \sigma^2)p(w'|\alpha)})$.

- To perform pure Gibbs sampling for w and \hat{t} we would like to first sample w from $p(w|\hat{t}, \hat{x}, t, x, \sigma^2, \alpha)$. As argued earlier, we can treat \hat{t} and \hat{x} as being part of t and x and so we might as well sample w from $p(w|t, x, \sigma^2, \alpha)$. Then we have that

$$\begin{aligned}
& p(w|t, x, \sigma^2, \alpha) \\
&= \frac{p(w, t, x, \sigma^2, \alpha)}{p(t, x, \sigma^2, \alpha)} && \text{by the product rule} \\
&\propto p(w, t, x, \sigma^2, \alpha) \\
&= \prod_{i=1}^N p(t_i|x_i, w\sigma^2)p(w|\alpha)
\end{aligned}$$

Hence, in Gibbs sampling, we can sample w proportionally to the joint distribution or specifically the distribution $\prod_{i=1}^N p(t_i|x_i, w\sigma^2)p(w|\alpha)$. As will be shown in the next bullet point, this distribution is the posterior of w and it is

$$w|t, x, \sigma^2, \alpha \sim \text{Normal}\left(\left[\frac{1}{\alpha}I + \frac{1}{\sigma^2}xx^T\right]^{-1} \frac{1}{\sigma^2}xt, \left[\frac{1}{\alpha}I + \frac{1}{\sigma^2}xx^T\right]^{-1}\right)$$

Next, we need to sample \hat{t} which requires sampling from $p(\hat{t}|\hat{x}, t, x, w, \sigma^2, \alpha)$. We can see that

$$\begin{aligned}
& p(\hat{t}|\hat{x}, t, x, w, \sigma^2, \alpha) \\
&= \frac{p(\hat{t}, \hat{x}, t, x, w, \sigma^2, \alpha)}{p(\hat{x}, t, x, w, \sigma^2, \alpha)} && \text{by the product rule} \\
&\propto p(\hat{t}, \hat{x}, t, x, w, \sigma^2, \alpha) \\
&= \prod_{i=1}^N p(t_i|x_i, w, \sigma^2)p(w|\alpha)p(\hat{t}|\hat{x}, w, \sigma^2) \\
&\propto p(\hat{t}|\hat{x}, w, \sigma^2)
\end{aligned}$$

Hence, Gibbs sampling, we can sample \hat{t} proportionally to the distribution $p(\hat{t}|\hat{x}, w, \sigma^2)$ which has distribution $\text{Normal}(w^T\hat{x}, \sigma^2)$.

- We compute the analytic form of the posterior predictive: $p(\hat{t}|\hat{x}, t, x, \sigma^2, \alpha)$. We begin by noting the following

$$\begin{aligned}
& p(\hat{t}|\hat{x}, t, x, \sigma^2, \alpha) \\
&= \int p(\hat{t}|\hat{x}, t, x, \sigma^2, w, \alpha) p(w|\hat{x}, t, x, \sigma^2, \alpha) dw && \text{by the sum rule} \\
&= \int p(\hat{t}|\hat{x}, t, x, \sigma^2, w, \alpha) p(w|\hat{x}, t, x, \sigma^2, \alpha) dw && \text{by the product rule} \\
&= \int p(\hat{t}|\hat{x}, \sigma^2, w) p(w|t, x, \sigma^2, \alpha) dw
\end{aligned}$$

Where the last line above follows by the fact that $p(\hat{t}|\hat{x}, t, x, \sigma^2, w, \alpha) = p(\hat{t}|\hat{x}, \sigma^2, w)$ and $p(w|\hat{x}, t, x, \sigma^2, \alpha) = p(w|t, x, \sigma^2, \alpha)$, the latter following by the fact that w can only have a dependency on \hat{x} if \hat{t} is known.

Next, we will argue that both $\hat{t}|\hat{x}, \sigma^2, w$ and $w|t, x, \sigma^2, \alpha$ are Gaussian random variables. Firstly, $\hat{t}|\hat{x}, \sigma^2, w \sim \text{Normal}(w^T \hat{x}, \sigma^2)$ by the likelihood given in the question. As for $w|t, x, \sigma^2, \alpha$ we note that

$$w|\alpha \sim \text{Normal}(0, \alpha I)$$

$$t|w, x, \sigma^2 \sim \text{Normal}(w^T x, \sigma^2)$$

Then we note that

$$p(w|t, x, \sigma^2, \alpha) \propto p(t|w, x, \sigma^2, \alpha)p(w, x|\sigma^2, \alpha) \propto p(t|w, x, \sigma^2)p(w|\alpha)$$

Then by 2.116, page 93, in Pattern Recognition and Machine Learning by Bishop it follows that

$$w|t, x, \sigma^2, \alpha \sim \text{Normal}\left(\left[\frac{1}{\alpha}I + \frac{1}{\sigma^2}xx^T\right]^{-1}\frac{1}{\sigma^2}xt, \left[\frac{1}{\alpha}I + \frac{1}{\sigma^2}xx^T\right]^{-1}\right)$$

Then by the earlier observations that

$$p(\hat{t}|\hat{x}, t, x, \sigma^2, \alpha) = \int p(\hat{t}|\hat{x}, \sigma^2, w)p(w|t, x, \sigma^2, \alpha)dw$$

and

$$\hat{t}|\hat{x}, \sigma^2, w \sim \text{Normal}(w^T \hat{x}, \sigma^2)$$

and by linear combination rules of Gaussian random variables we have that

$$\hat{t}|\hat{x}, t, x, \sigma^2, \alpha \sim \text{Normal}\left(\left[\left(\frac{1}{\alpha}I + \frac{1}{\sigma^2}xx^T\right)^{-1}\frac{1}{\sigma^2}xt\right]^T \hat{x}, \hat{x}^T \left(\frac{1}{\alpha}I + \frac{1}{\sigma^2}xx^T\right)^{-1} \hat{x} + \sigma^2\right)$$

This completes the derivation of the analytic form of the posterior predictive.

5.